

PROGRAMMING and ALGORITHMS II



INTRODUCTION TO
**ADVERSARIAL
SEARCH**

MINIMAX AND ALPHA-BETA REDUCTION

Algorithm Design Paradigms

DIVIDE & CONQUER

Break down a problem into independent sub-problems of related type, solve them separately and combine the solutions.

GREEDY APPROACH

Use a sequence of locally optimal decisions incrementally to build up a solution.

DYNAMIC PROGRAMMING

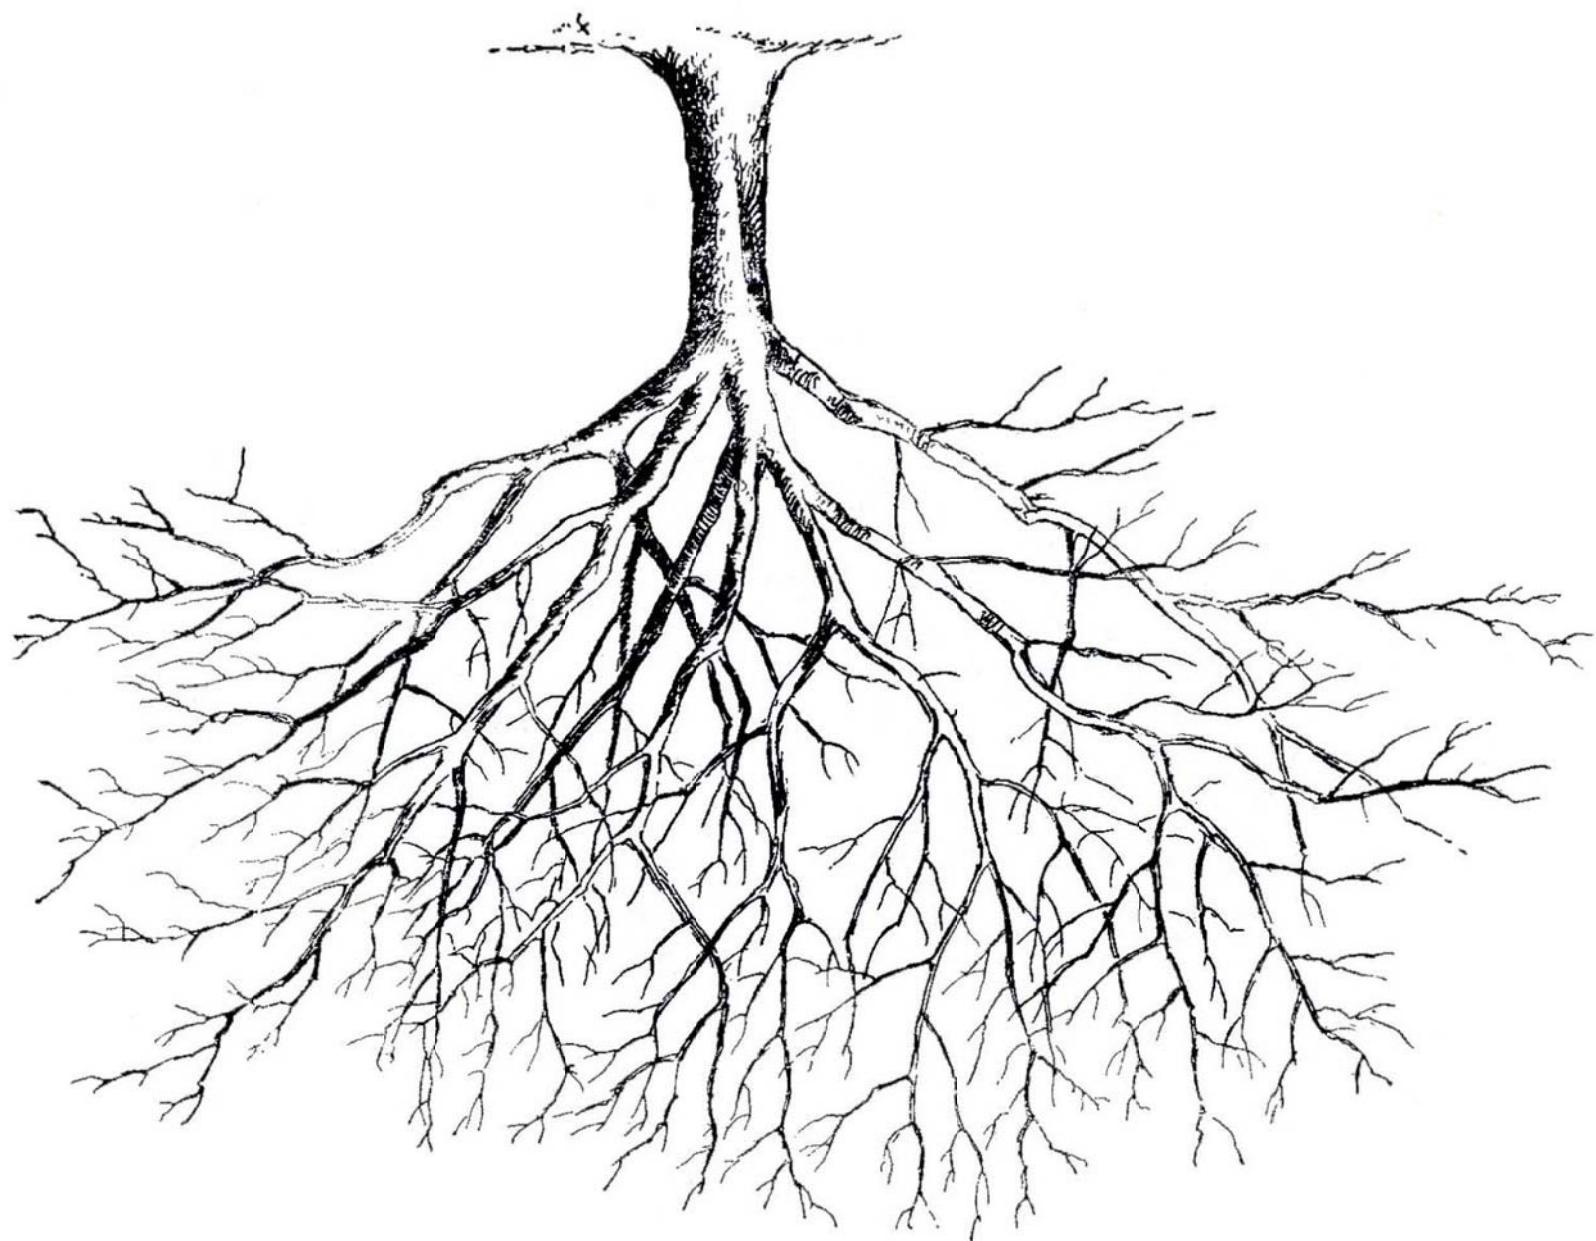
Break down a problem into overlapping sub-problems of related type, build up solutions from larger and larger sub-solutions.

Deep Blue

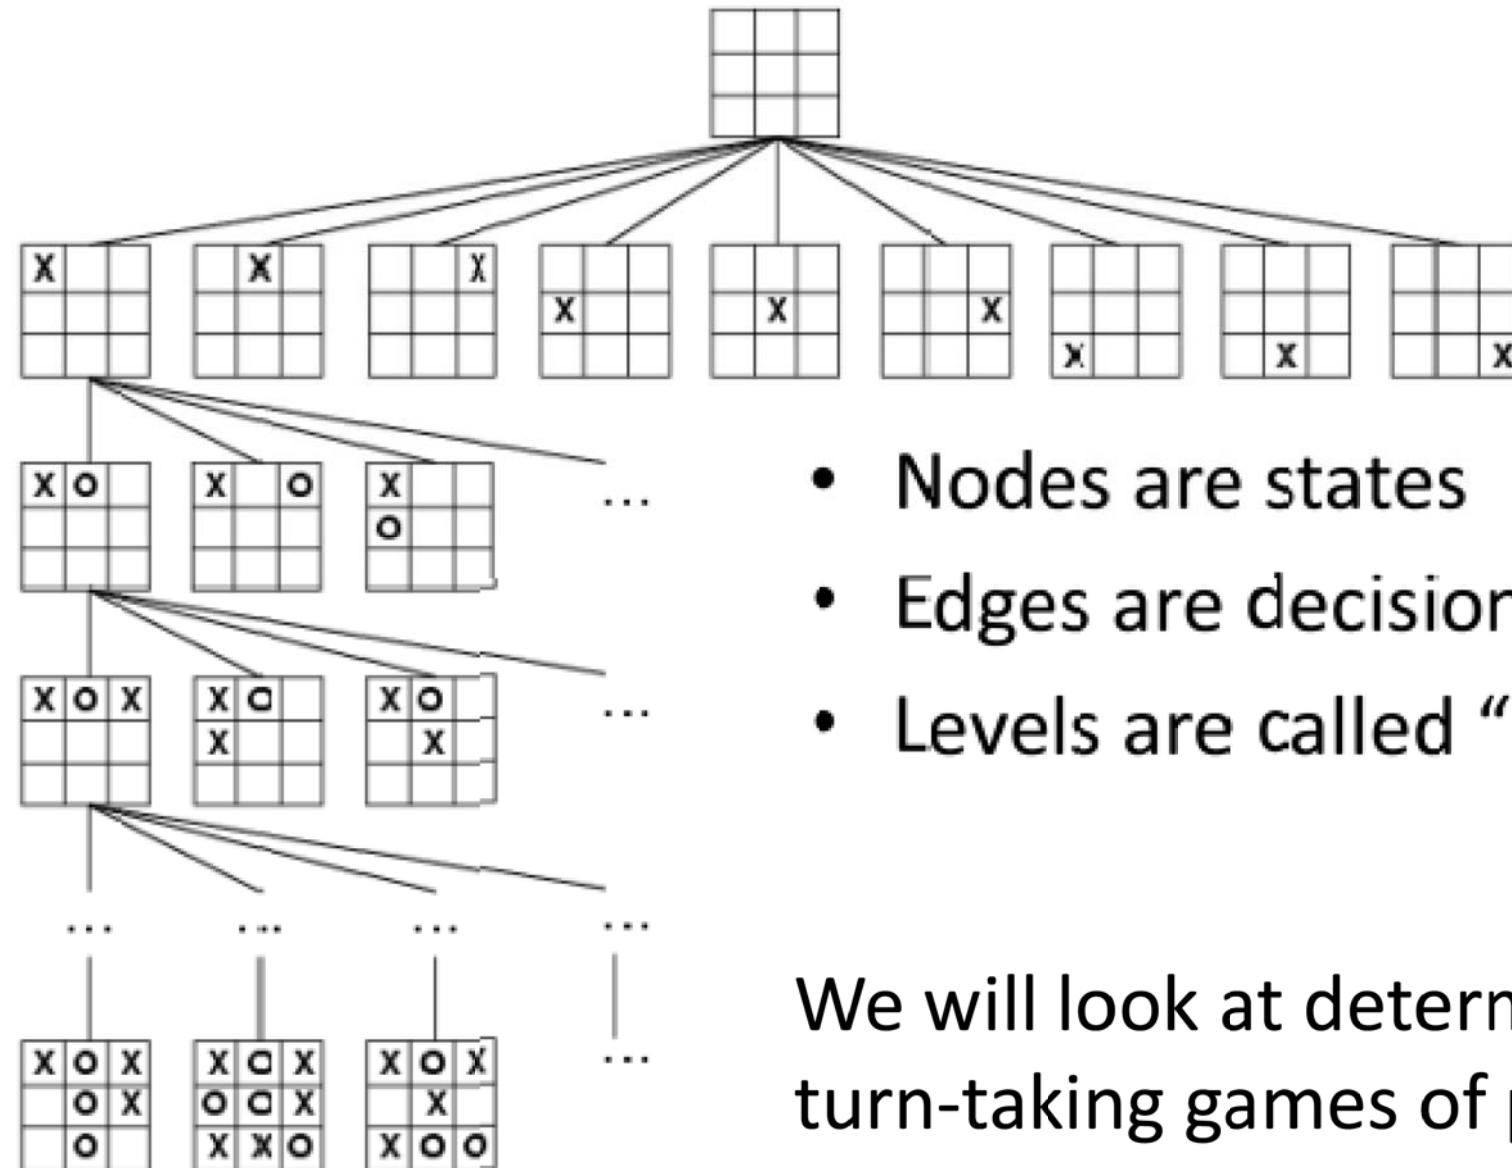
- In 1996 Deep Blue became the first machine to win a chess game against a reigning world champion (Garry Kasparov) under regular time controls
- 30 nodes running at 120 MHz
- program written in C exploring a ‘game tree’ at 200 million positions per second



Concept of a ‘Game Tree’



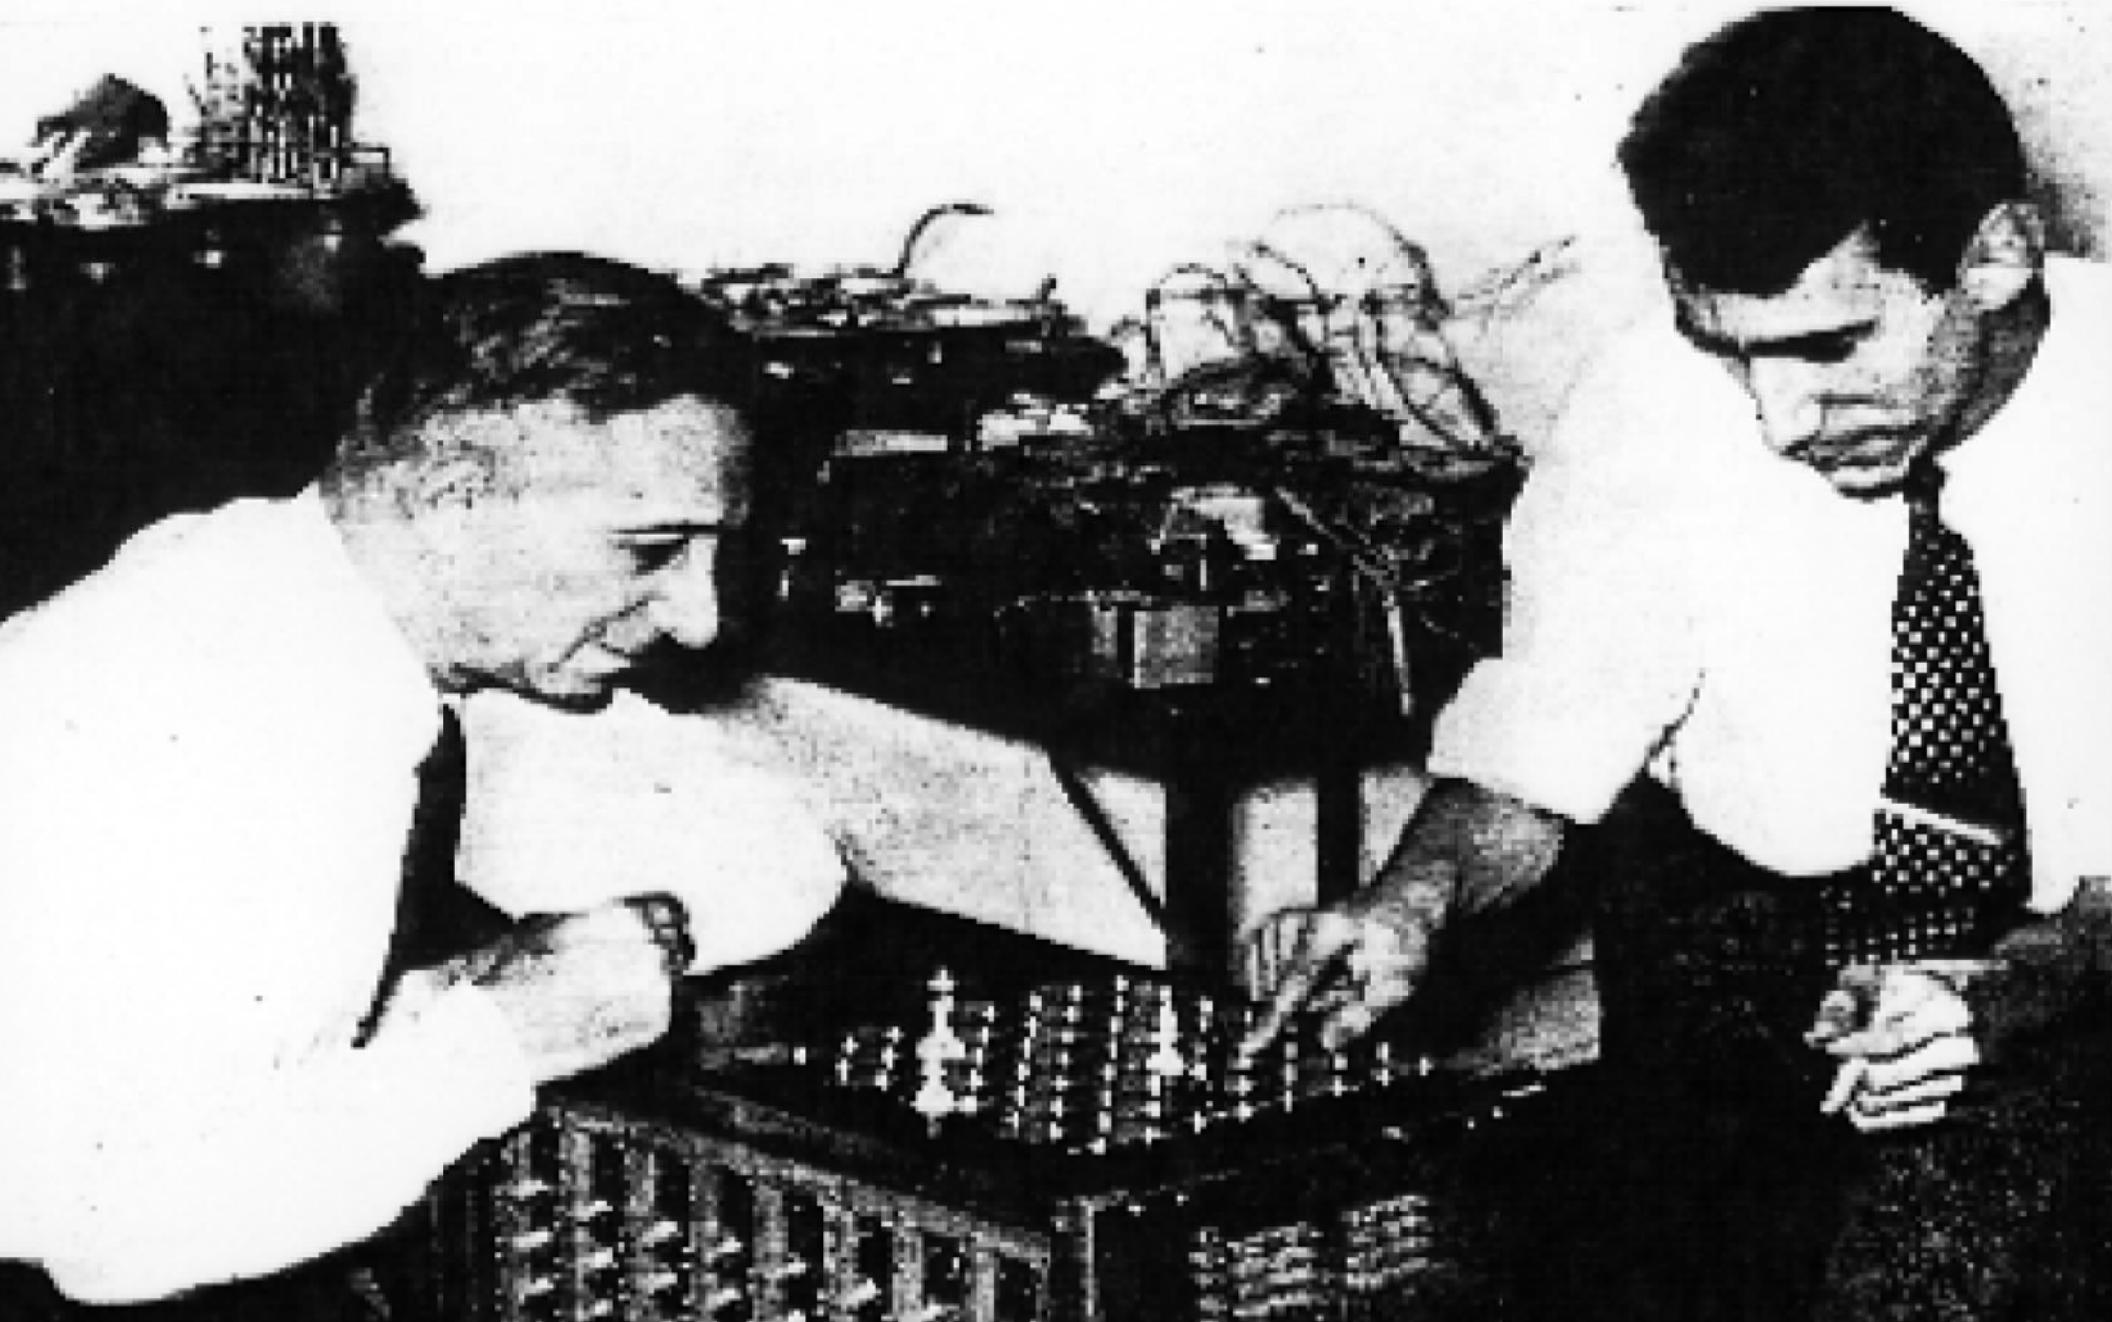
Semantics of a ‘Game Tree’



Shannon's Chess Machine



Department of
Computer Science



- In 1950 Shannon published paper:
Programming a Computer for Playing Chess.
- General Idea: Computer simulates sequences of the enemy as well as own moves.
- The process for having the computer decide on which move to make is guided by trying to Minimize the maximum possible loss, or Maximize the minimum win...
→ Find winning sequence where enemy plays optimally, e.g. maximizing win utility which is $+1$ for win, 0 for tie, and -1 for loss

Game Tree of OXO

computer's
turn

MAX (X)

opponent's
turn

MIN (O)

computer's
turn

MAX (X)

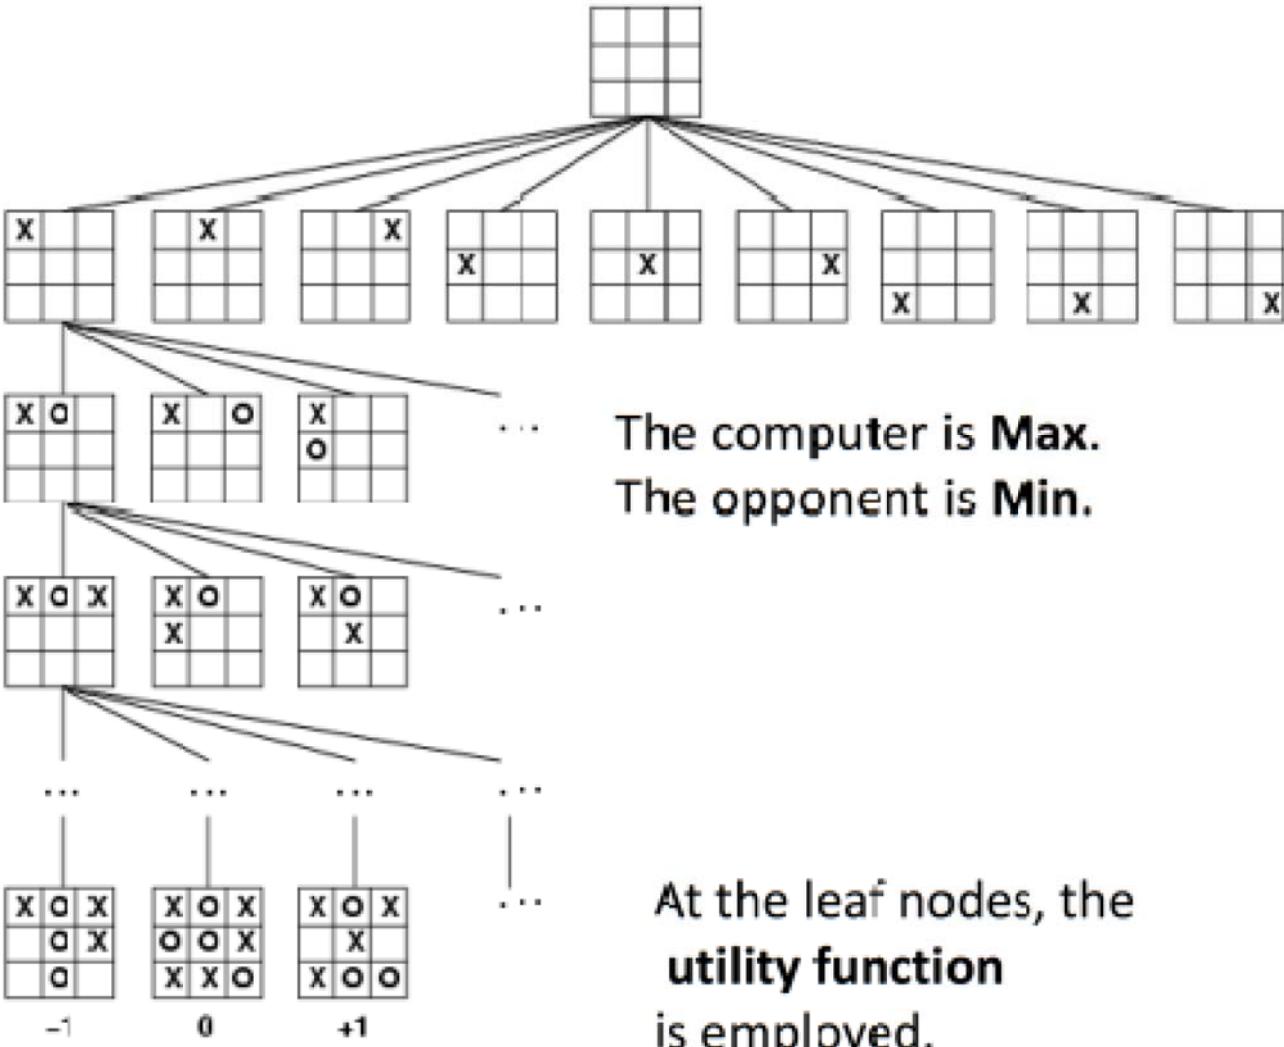
opponent's
turn

MIN (O)

leaf nodes
are evaluated

TERMINAL

Utility



The MiniMax Algorithm

function MINIMAX-DECISION(*state*) **returns** *an action*

v \leftarrow MAX-VALUE(*state*)

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

v $\leftarrow -\infty$

for *s* in SUCCESSORS(*state*) **do**

v \leftarrow MAX(*v*, MIN-VALUE(*s*))

return *v*

function MIN-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

v $\leftarrow \infty$

for *s* in SUCCESSORS(*state*) **do**

v \leftarrow MIN(*v*, MAX-VALUE(*s*))

return *v*

- Optimality: Yes (against an optimal opponent)
- Completeness: Yes (if tree is finite)
- Time complexity: $O(m^n)$
- Space complexity: $O(mn)$ (depth-first exploration)

Example: Chess ... $m \approx 30$ options per move, $n \approx 100$ moves for "reasonable" games
→ exact solution completely infeasible

- Assume we can evaluate 10,000 boards per second and we have 10 secs → 10^5 nodes per move

Example: Chess.... $m^n = 10^5$, $m = 30 \rightarrow n \approx 3$

3-ply look-ahead is a truly hopeless chess player!

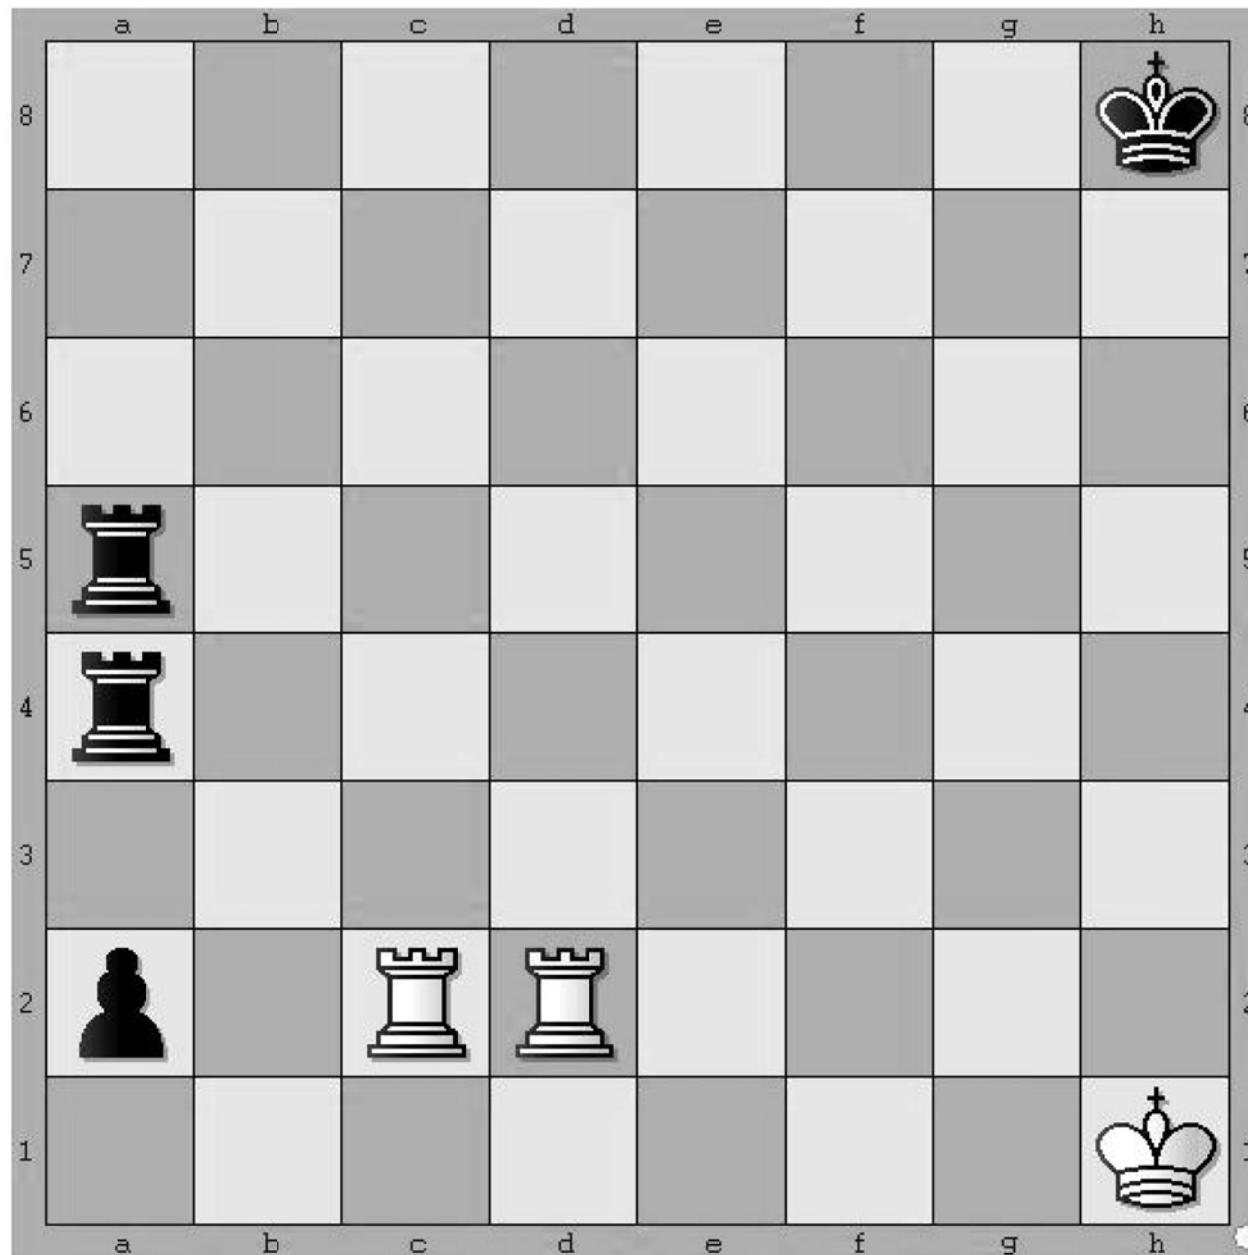
- 4-ply ≈ human novice
- 8-ply ≈ typical PC, human master
- 12-ply+ ≈ Deep Blue, Kasparov

(Remember: A computer which evaluates its own legal moves plus the legal responses to those moves is searching to a depth of 2-ply...)

Shannon gave a rough example of an evaluation function for chess in which the value of the black position was subtracted from that of the white position.

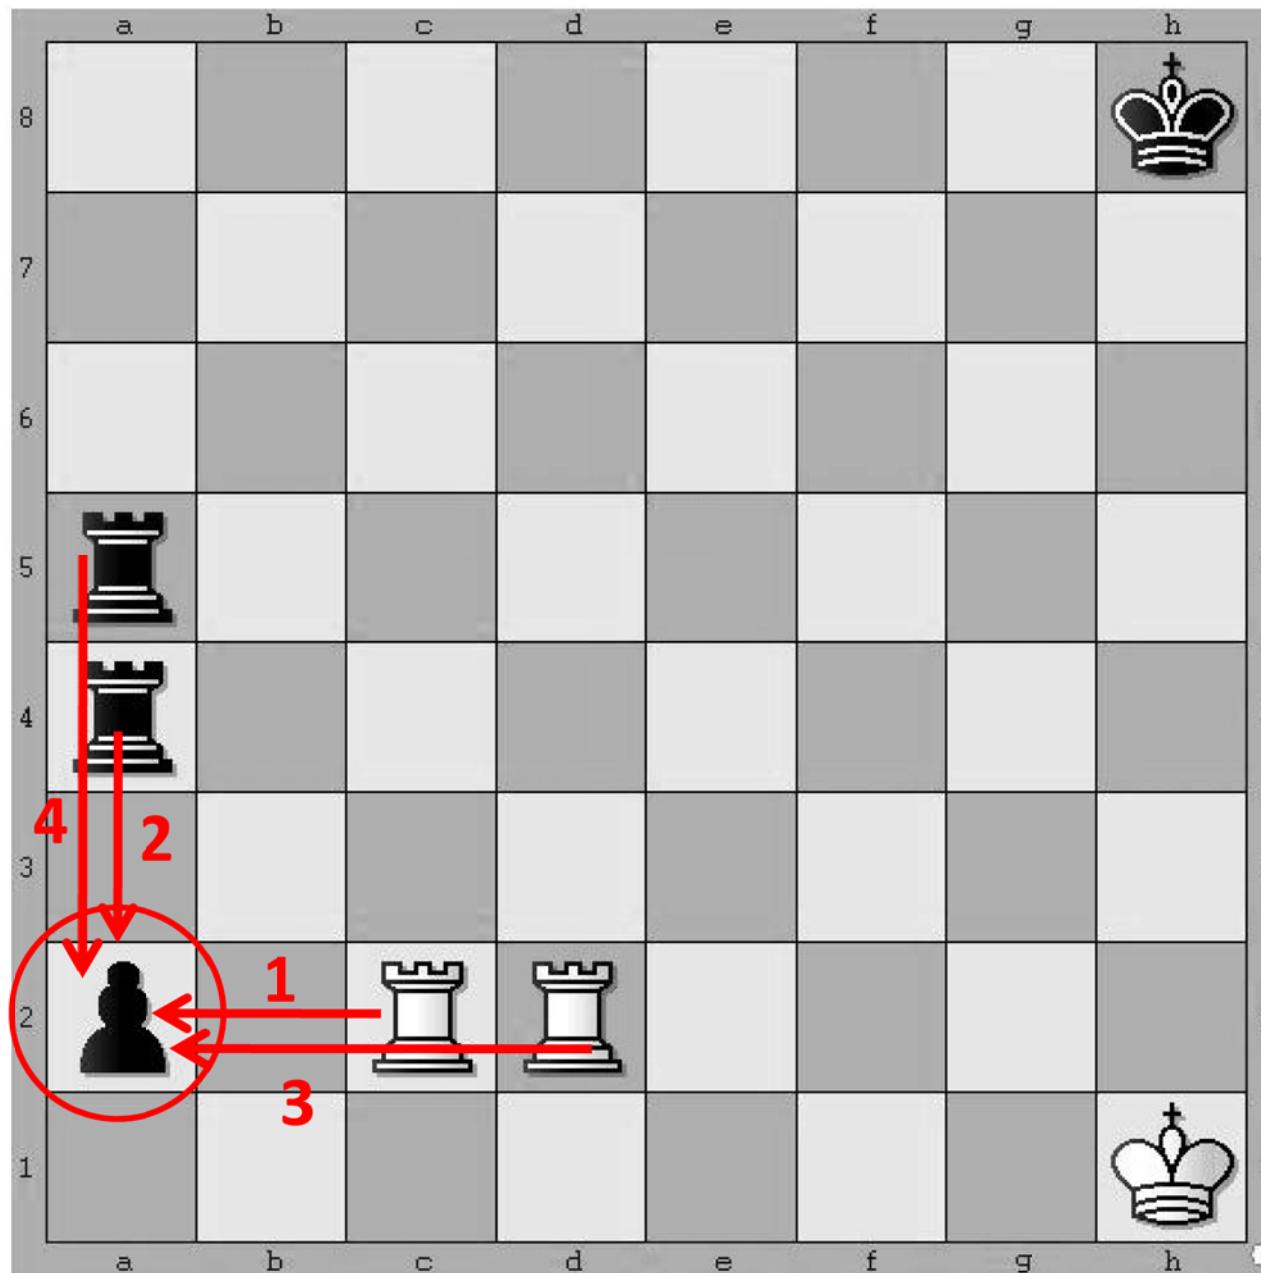
- **Piece Value:** was counted according to the relative chess piece value (1 point for a pawn, 3 points for a knight or bishop, 5 points for a rook, and 9 points for a queen).
- **Positional Factors:** subtracting $\frac{1}{2}$ point for each doubled pawns, backward pawn, isolated pawn.
- **Mobility:** adding 0.1 point for each legal move available.
- **Checkmate:** king has artificial value of 200 points

Example: Look-ahead 3



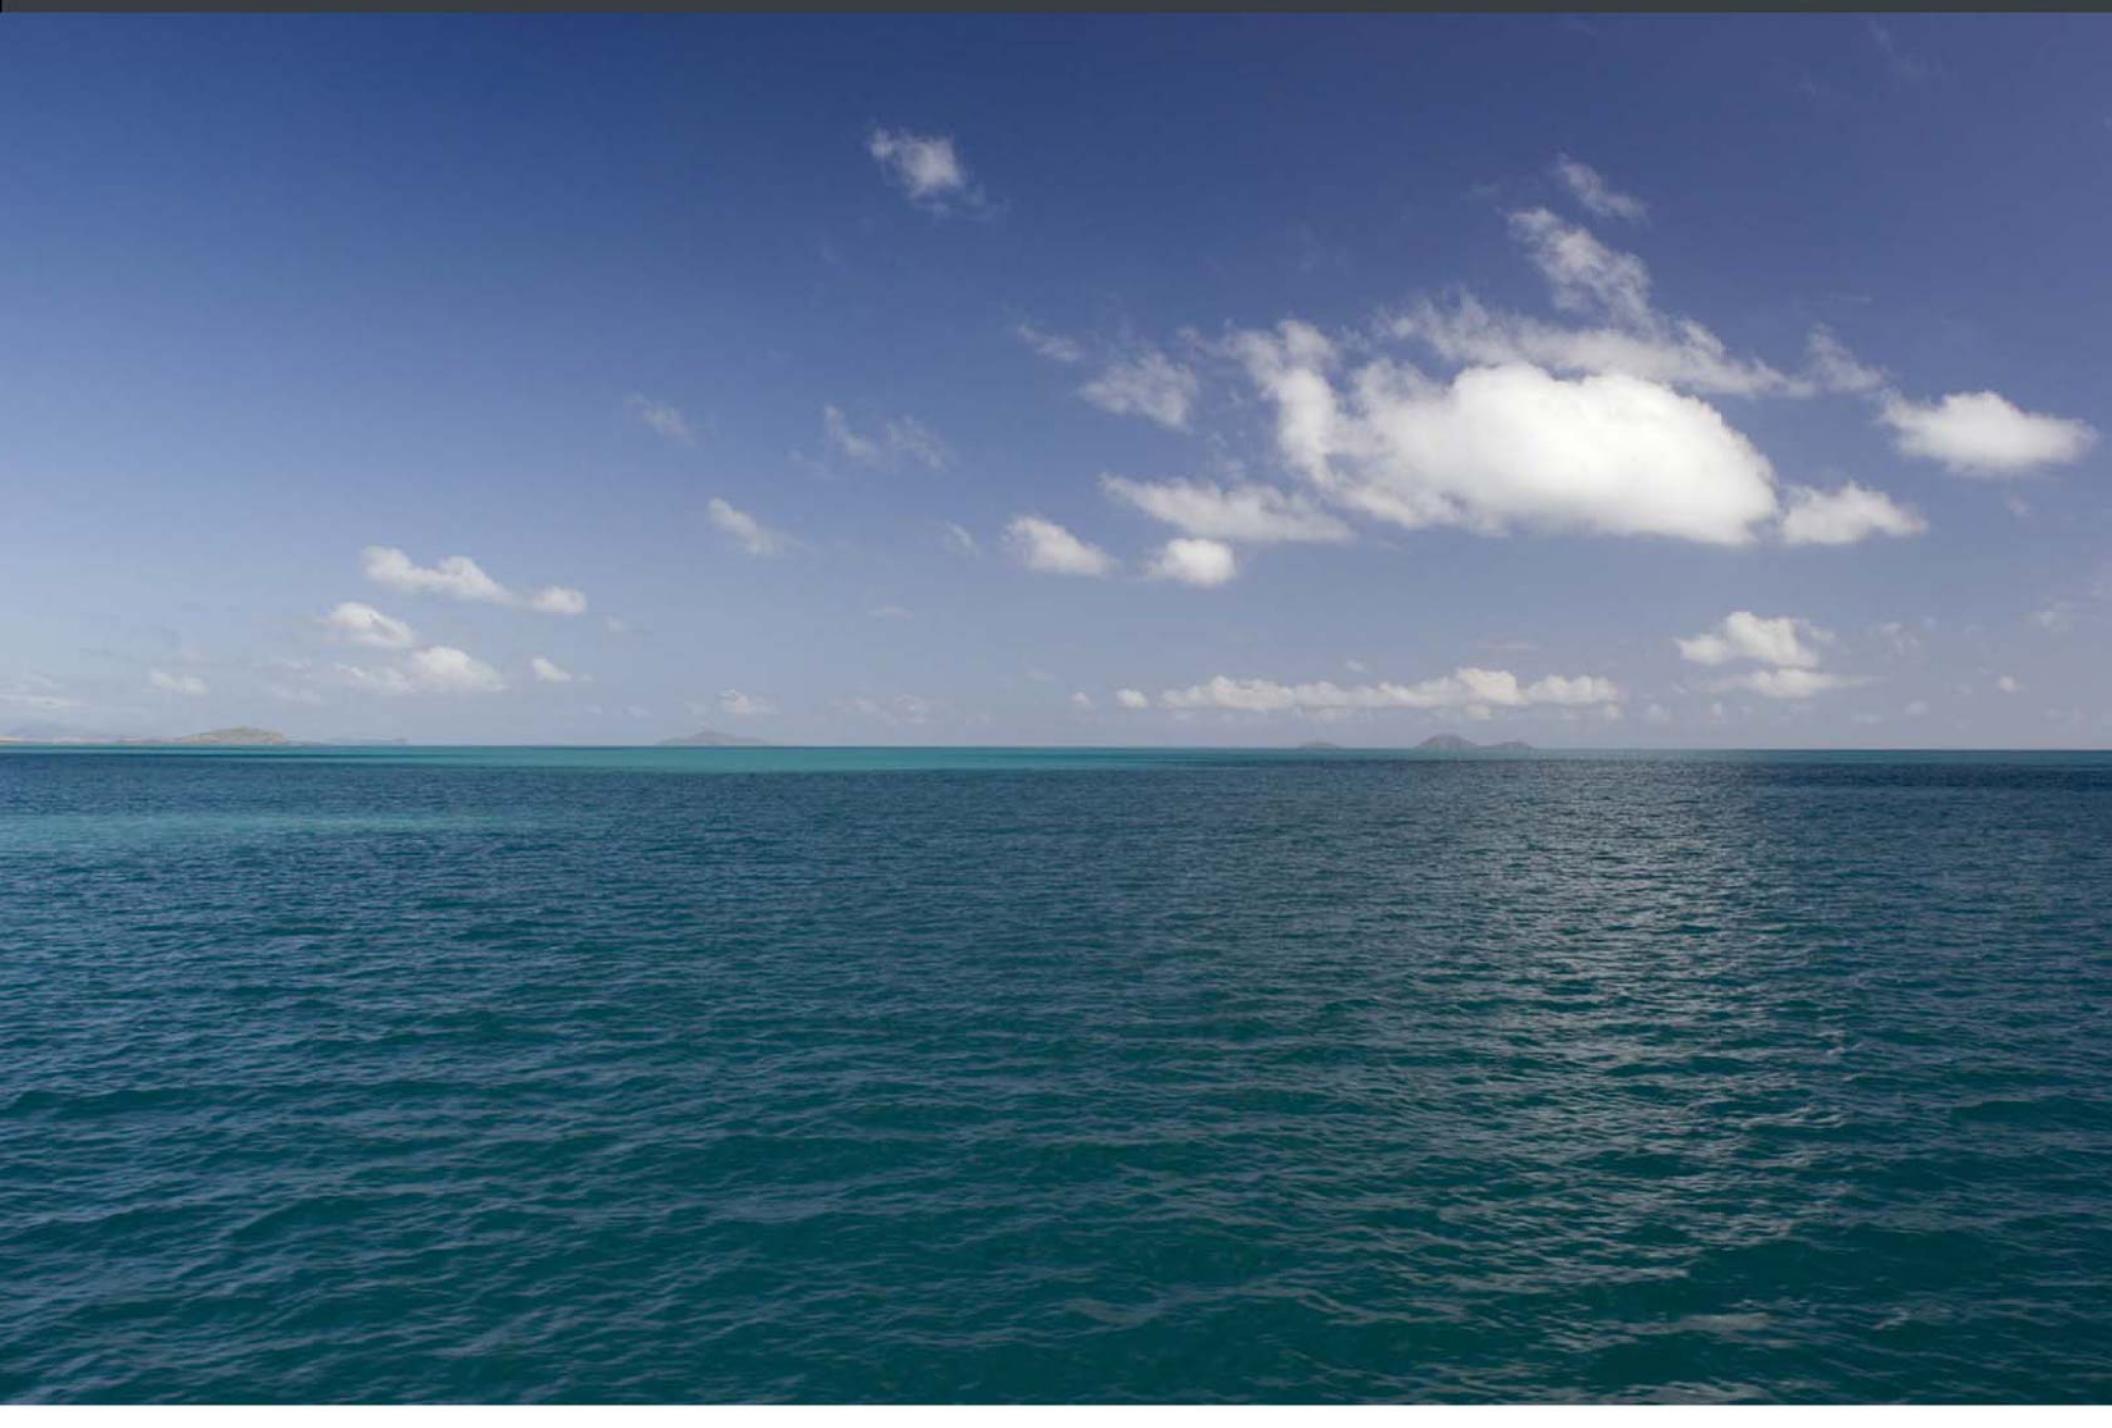
Example by Jonathan Pettersson

Horizon Effect Example



Example by Jonathan Pettersson

Horizon Effect



The Horizon Effect

- Your algorithm searches to depth n
- What happens if:
 - Evaluation(s) at depth n is very positive
 - Evaluation(s) at depth $n+1$ is very negative
- Or:
 - Evaluation(s) at depth n is very negative
 - Evaluation(s) at depth $n+1$ is very positive

Will this ever happen
in practice?



- Often it is useful to look deeper into game tree
→ peak past the horizon...
- What nodes to explore?
- Human players have some intuition about move quality
 - “Interesting” vs “boring”
 - “Promising” vs “dead end”
- Expand horizon for potential high impact moves
- Quiescence Search finds strong changes OR consistency along move sequences

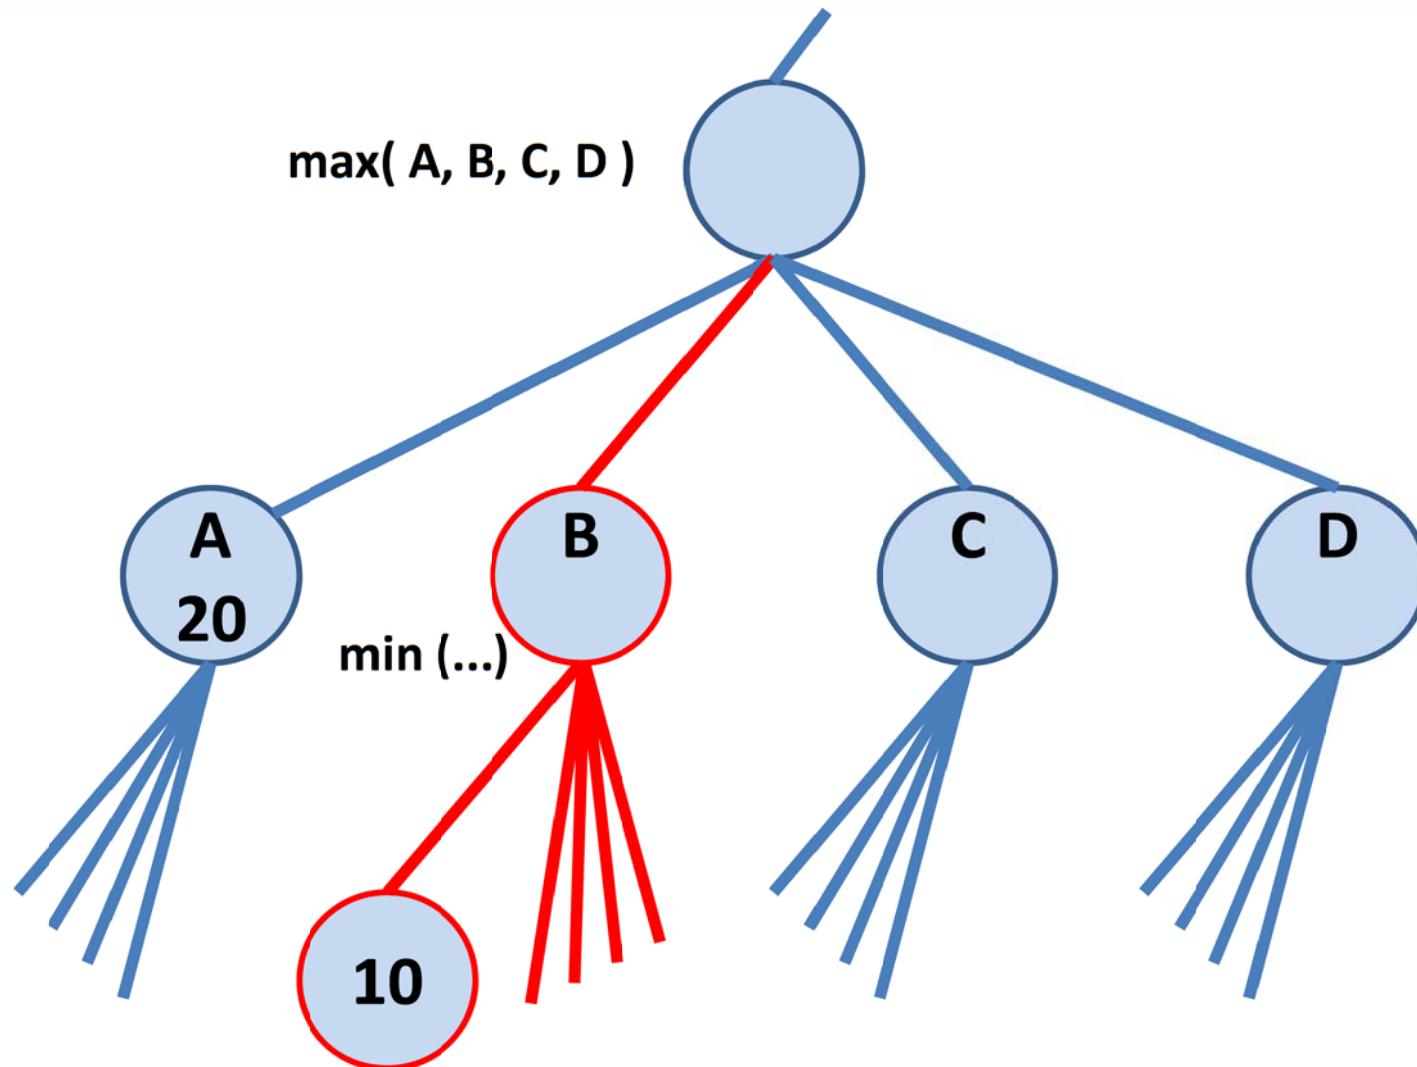
- quiescent search does not have a set depth
- a bad ordering of the moves may result in extremely deep lines of search
 - this is called a *quiescent search explosion*
- ordering the search space can provide a limiting criterion, e.g. MVV/LVA (Most Valuable Victim/Least Valuable Attacker)
- this way we are not calculating obvious losses early in the search
- However, horizon effect is still present: we can not know if captured pieces are protected. (A queen capturing an unprotected pawn may be a good move, but a majority of the time the pieces will be protected)

- **Opening Books:** pre-compute opening move sequences and store them for direct access (i.e. pre-analyse root region of game tree)
- **Endgame Databases:** store endgame patterns and exhaustively pre-compute endgame (i.e. pre-analyse leaf regions of game tree)
- **Killer Move Analysis:** store and re-identify configurations likely to lead to endgame (i.e. identify critical/forced paths to end game)
- **Transposition (Dynamic Programming) Tables:** different move sequences may result in identical boards, computing solutions once and storing analysis results can avoid double exploration

Pruning Game Trees

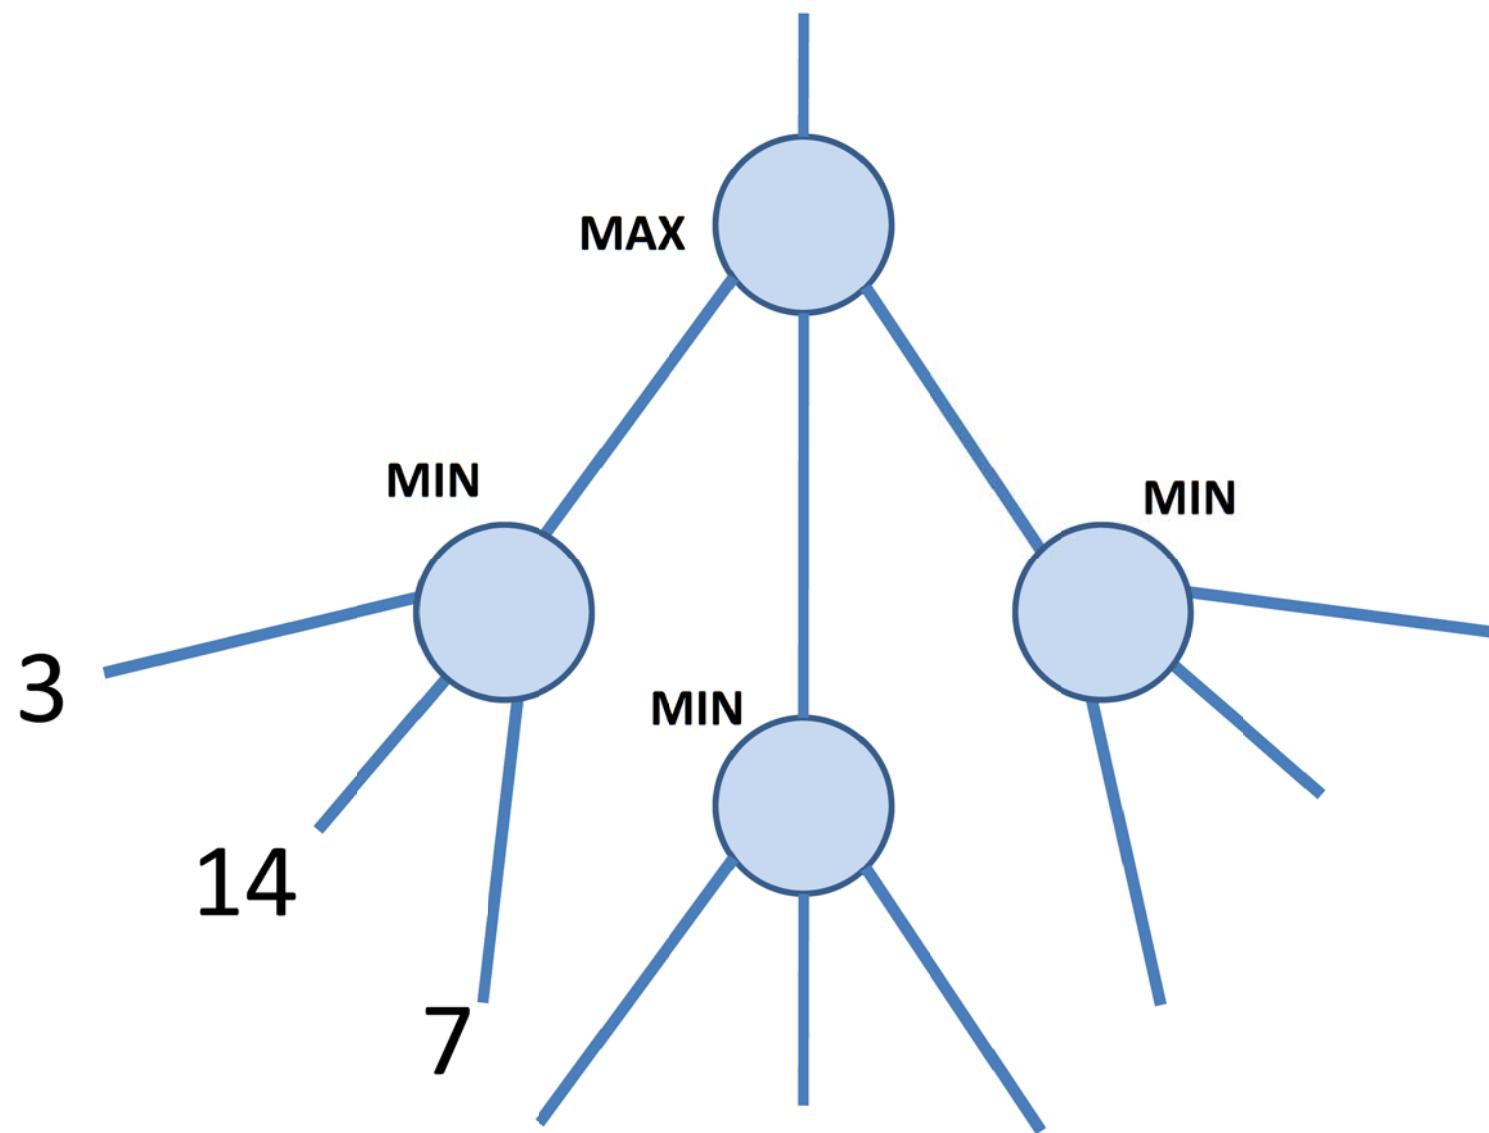


Idea of Game Tree Pruning

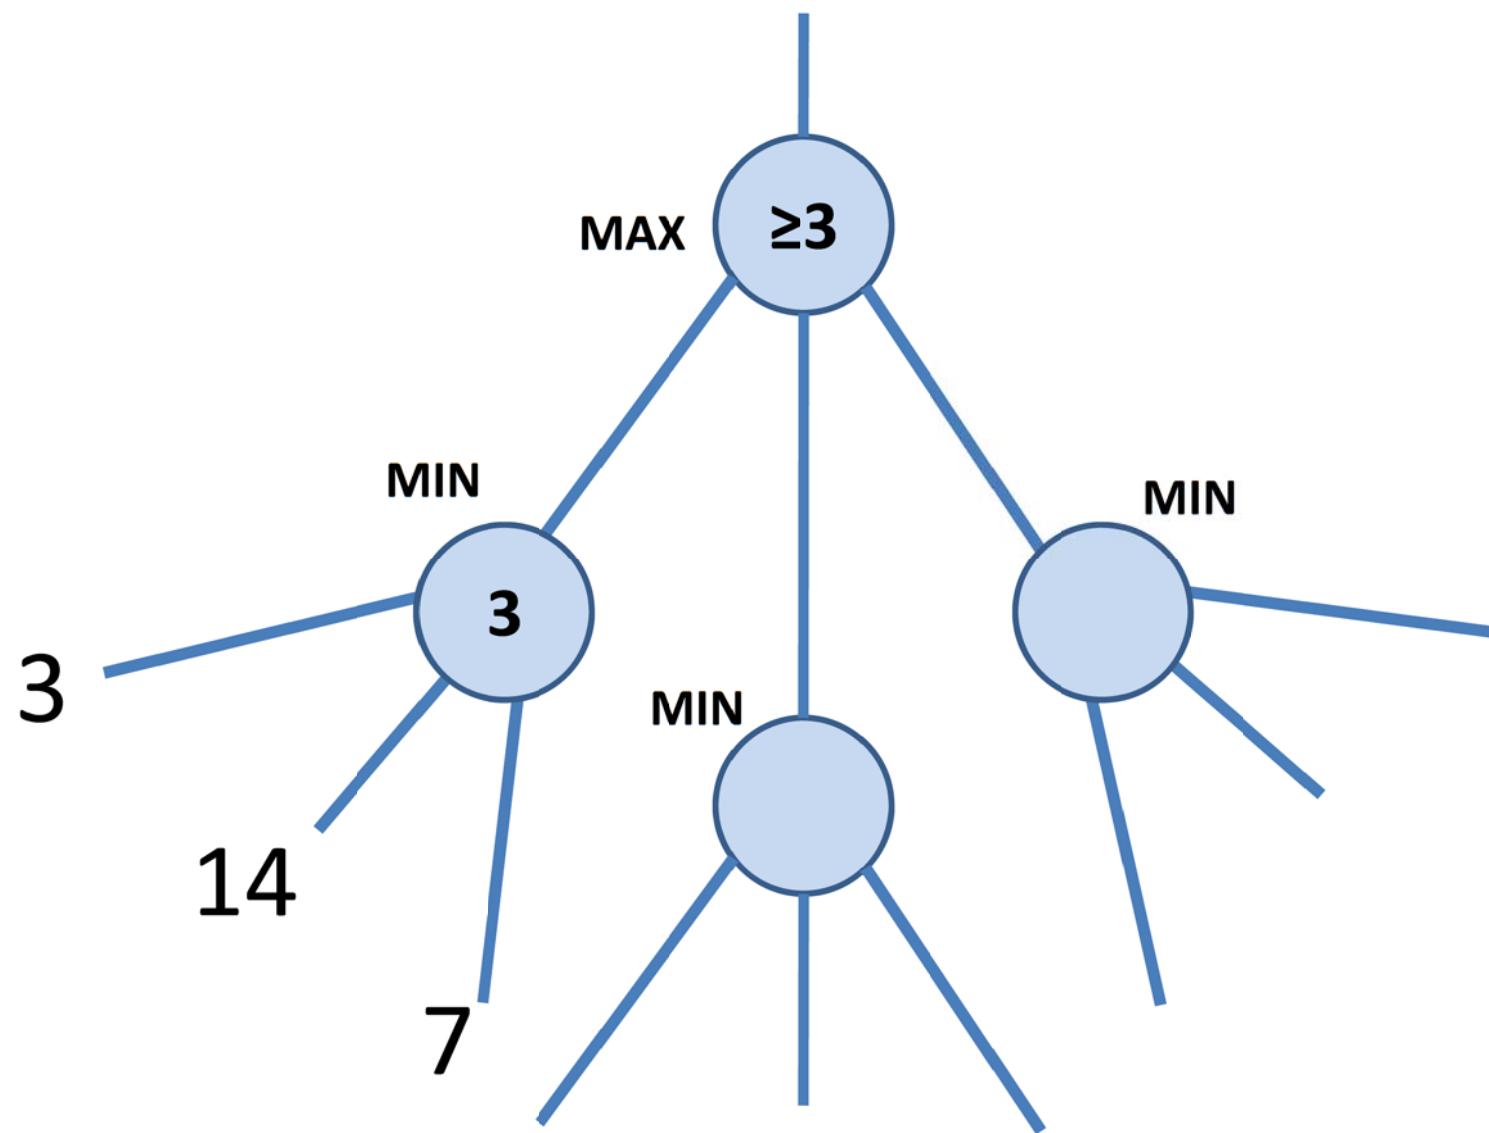


Node B and its entire subtree have no influence on the value of parent node, since $\max(20, \dots) \geq 20$ and $B \leq 10 \dots$

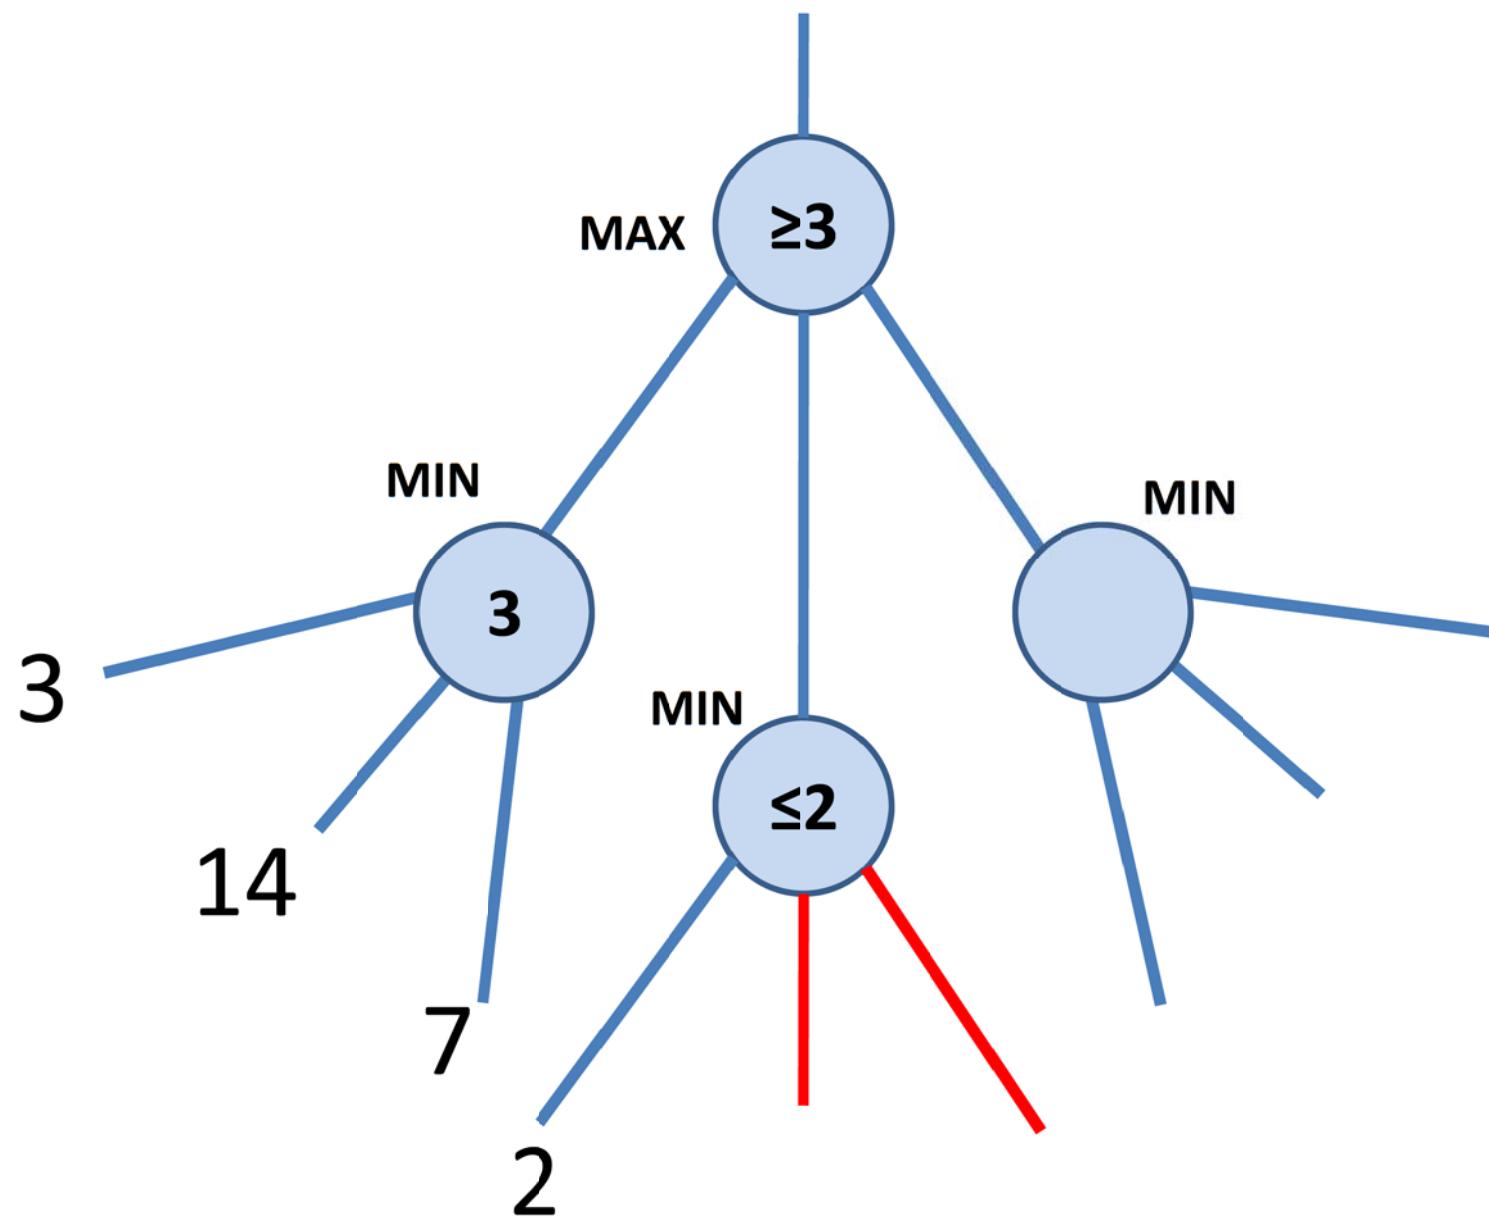
Alpha-Beta Pruning Example



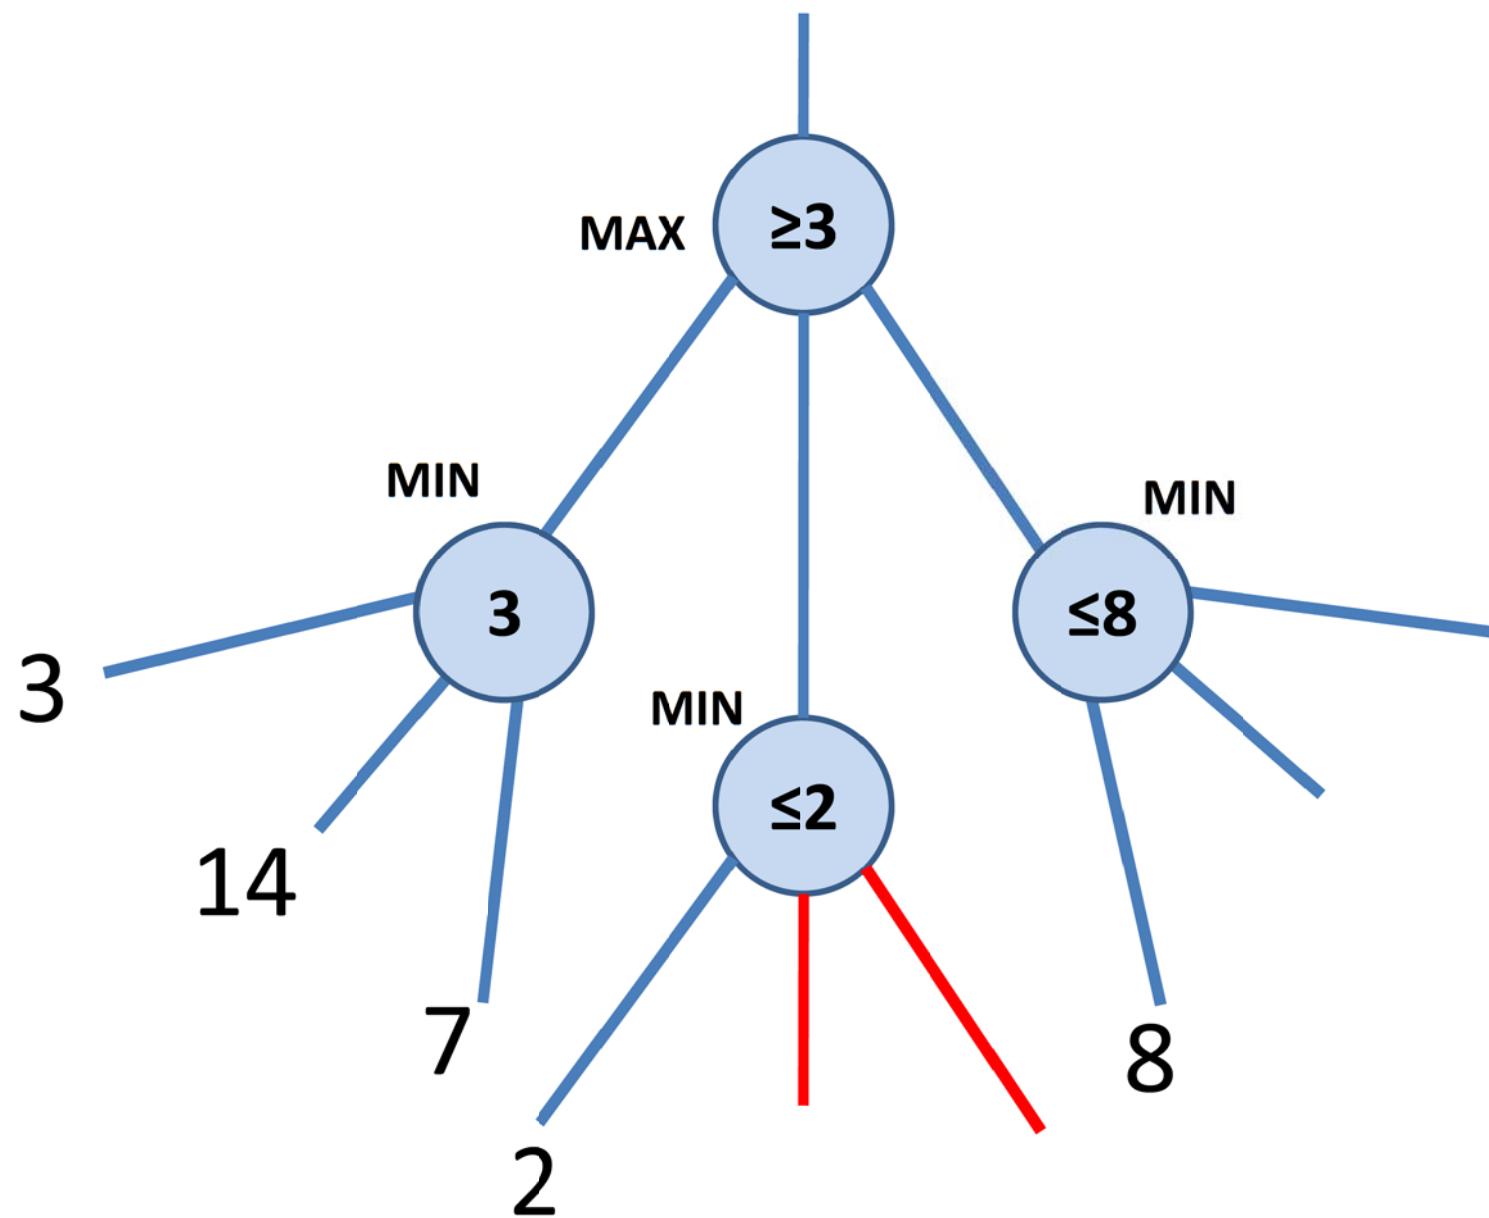
Alpha-Beta Pruning Example



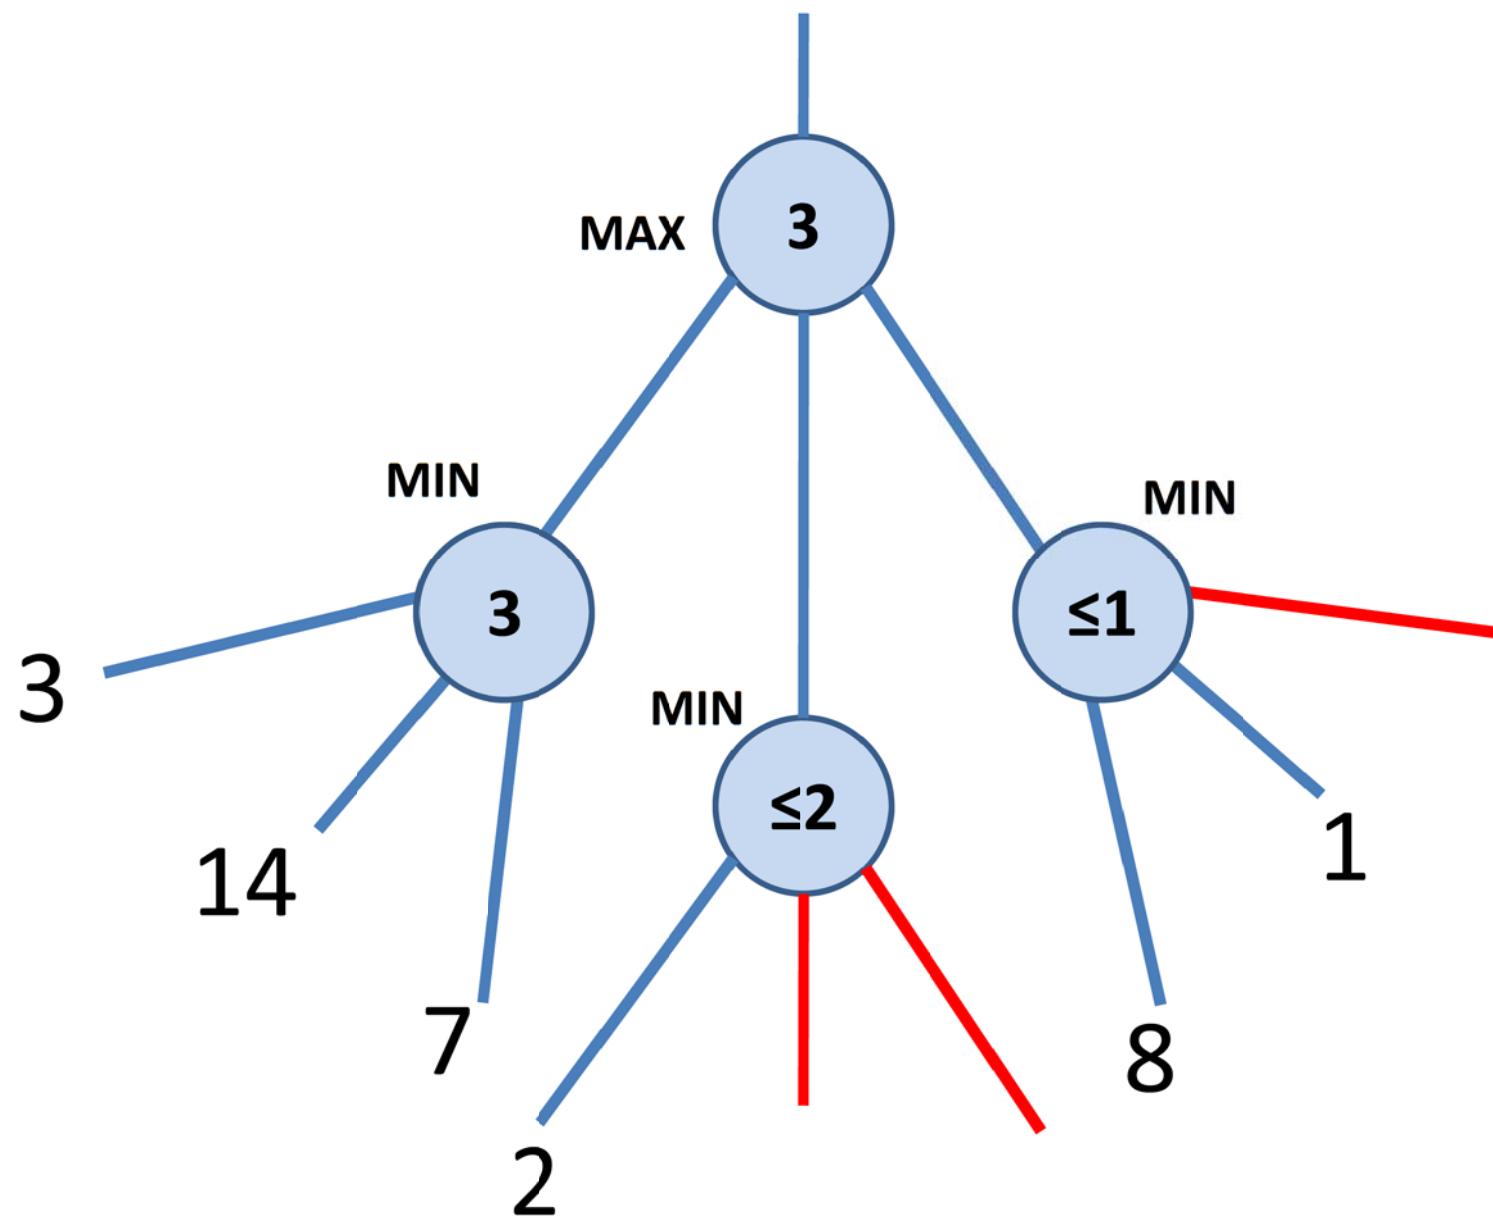
Alpha-Beta Pruning Example



Alpha-Beta Pruning Example



Alpha-Beta Pruning Example



- Alpha-Beta is computes the same value for the root node as computed by the MiniMax algorithm
- **Worst Case:** Alpha-Beta does NO pruning, examining m^n leaf nodes, where nodes have m children and a n -ply search is performed
- **Best Case:** Alpha-Beta will examine only $2m^{n/2}$ leaf nodes. Ergo, the search can now be twice as deep as MiniMax...
- If possible, consider ordering of leaf nodes to optimise pruning performance...

Alpha-Beta Algorithm

function ALPHA-BETA-SEARCH(*state*) **returns** *an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\textit{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) **returns** *a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *t* : in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

function MIN-VALUE(*state*, α , β) **returns** *a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*

$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*

Deep Blue vs Kasparov



Draughts Championship



Department of
Computer Science

Silicon Graphics Draughts Championship



Silicon Graphics
World Draughts Championship

GO is still Human-reigned



- Game playing can be effectively modeled as a **Search Problem**
- **Game Trees** represent alternate computer/opponent moves
- Evaluation functions estimate the quality of a given board configuration
- **The MiniMax Algorithm** is a procedure which chooses moves by assuming that the opponent will always choose the move which is best for them
- **Alpha-Beta Pruning** is a procedure which can discard large parts of the search tree and, thus, allow search to go deeper