

Trees



Searching

3

Suppose we want to search for things in a list

One possibility is to keep the items in a 'randomly' ordered list, so insertion is $O(1)$, but then a search takes $O(n)$ time

Or, we could keep them in a sorted list, in which case we can use a binary search which takes $O(\log(n))$ time, but then new items would have to be added in the middle, which takes $O(n)$ time

When there is a mixture of search and insert operations, and both operations need to be well below $O(n)$, then the items can usefully be stored in an ordered binary tree

A tree is created out of cells, with each cell having two pointers

We will create ordered binary trees, without worrying about how well balanced the tree is

Balancing techniques include:

- Reorder or randomise the input data (assuming few updates)
- Occasional re-construction of the tree
- Use a self-balancing (red-black, AVL, 2-3, ...) tree

Tree structure

6

Here's a struct for holding one node in a tree of ints:

```
struct node {  
    struct node *left;  
    int key;  
    struct node *right;  
};  
typedef struct node node;
```

This is essentially the same as

`Tree a = Tip | Node (Tree a) a (Tree a)`
in Haskell (using `NULL` for `Tip`)

New node

7

Here's a function to create a new node (a one-element tree):

```
node *new_node(int n) {  
    node *p = malloc(sizeof(node));  
    *p = (node) { NULL, n, NULL };  
    return p;  
}
```

Recursive Insertion

Here's a recursive insertion function:


```
node *insert_node(node *p, int n) {  
    if (p == NULL) p = new_node(n);  
    else if (n < p->key)  
        p->left = insert_node(p->left, n);  
    else if (n > p->key)  
        p->right = insert_node(p->right, n);  
    return p;  
}
```

It uses `p` as a current-node variable

When you call it, it returns a possibly updated node, which you have to put back where you got it

Visualisation

9



P = NULL

Alternative Recursive Insertion 10

Here's a version which doesn't return anything, but uses a pointer to a pointer:


```
void insert_node(node **p, int n) {  
    if (*p == NULL) *p = new_node(n);  
    else if (n < (*p)->key)  
        insert_node(&(*p)->left, n);  
    else if (n > (*p)->key)  
        insert_node(&(*p)->right, n);  
}
```

It updates in place, and only does it once

When we reach the right place, we have a *pointer* to NULL, so we can replace it

Visualisation

11



Iterative Insertion

12

Here's a complicated iterative version:

```
node *insert_node(node *p, int n) {
    bool done = false;
    if (p == NULL) { p = new_node(n); done = true; }
    while (!done) {
        if (n == p->key) done = true;
        else if (n < p->key) {
            if (p->left != NULL) p = p->left;
            else { p->left = new_node(n); done = true; }
        }
        else {
            if (p->right != NULL) p = p->right;
            else { p->right = new_node(n); done = true; }
        }
    }
    return p;
}
```

Alternative Iterative Insertion

13

The structure is simpler using a pointer to a pointer:

```
void insert_node(node **p, int n) {  
    bool done = false;  
    while (!done) {  
        if (*p == NULL) {  
            *p = new_node(n); done = true;  
        }  
        else if (n == (*p)->key) done = true;  
        else if (n < (*p)->key) p = &(*p)->left;  
        else p = &(*p)->right;  
    }  
}
```

Choosing

14

Should you use recursive or iterative insertion, and should you use pointers-to-pointers or not?

You can disregard what anybody says about efficiency – what matters is balance and complexity issues

Use whichever you like – but when you write functions which use *both* left and right subtrees instead of just one, recursion stays simple while iteration gets nastier

A wrapper

15

Functions on trees are inconvenient if we force callers to use the nodes directly, either they have to catch the output, or pass a pointer-to-a-pointer

So we need a wrapping structure for a tree:

```
struct tree {  
    struct node *root;  
};  
typedef struct tree tree;
```

This can also be a useful place to put global information about the tree

New tree

16

Here's a reasonable function to create a new tree:

```
tree *new_tree() {  
    tree *t = malloc(sizeof(tree));  
    t->root = NULL;  
    return t;  
}
```


A wrapped insertion function might be:

```
void insert(tree *t, int n) {  
    t->root = insert_node(t->root, n);  
}  
// ----- OR -----  
void insert(tree *t, int n) {  
    insert_node(&t->root, n);  
}
```

In the iterative cases, the `insert` wrapper function can be combined with `insert_node` to form a single function

Recursive searching

18

Searching is a bit simpler, and can also be done recursively or iteratively – here's a recursive version:

```
node *find_node(node *p, int n) {  
    if (p == NULL) { }  
    else if (n < p->key)  
        p = find_node(p->left, n);  
    else if (n > p->key)  
        p = find_node(p->right, n);  
    return p;  
}
```

Iterative searching

19

Here's an iterative version:

```
node *find_node(node *p, int n) {  
    bool done = false;  
    while (!done) {  
        if (p == NULL) done = true;  
        else if (n == p->key) done = true;  
        else if (n < p->key) p = p->left;  
        else p = p->right;  
    }  
    return p;  
}
```

Wrapping

20

Again, we want a wrapper function

It shouldn't export any nodes to the user, it should just return (e.g.) a boolean to say whether the number is in the tree or not

```
bool contains(tree *t, int n) {  
    return find_node(t, n) != NULL;  
}
```

A *map* is a structure which maps keys to values

For example, when counting words, you might want to map word strings as keys, to integer counts as values:

```
struct node {  
    struct node *left;  
    char word[20];  
    int count;  
    struct node *right;  
};
```

The tree would be structured according to the words, and functions would retrieve or update the counts

Self balancing trees

22

There are many types of self-balancing tree, with red-black trees being the most popular in libraries because you only need one extra bit per node


The different types (AVL trees, 2-3 trees, ...) all use the same *mechanism* but have different *policies*

As Conor says "they are all morally the same"

Rotation

23

The mechanism used for balancing is rotation:



A rotate function

24

Here's a function to rotate right:

```
node *rotate_right(node *q) {  
    node *p = q->left;  
    q->left = p->right;  
    p->right = q;  
    return p;  
}
```

A *policy* is an algorithm which decides what rotations to do and when, according to some extra info in each node, and which guarantees $O(\log(n))$ depth